GREGORY JEFFREY WILLS

443-883-5856 • gregoryjwills11@gmail.com • https://github.com/gjw13

BACKGROUND AND EXPERIENCE SUMMARY

Seasoned DevOps engineer with over six years of experience solving challenging problems designing, developing, and optimizing high-performance parallel systems and data infrastructure. Proven expertise with distributed event streaming platform technologies and concepts, transforming siloed operations to unified, real-time data pipelines that drive informed decision-making.

WORK EXPERIENCE

Software Engineer, Viasat – Boston, MA **September 2021 – Present**

- Streamlined data stream provisioning by building a custom Terraform provider in Go, unifying workflows for 5+ engineering teams
- Accelerated service deployment by developing a k8s ExternalDNS webhook, slashing DNS provisioning time from hours to seconds
- Stabilized high-volume Kafka clusters (50B+ events/day) by deploying Cruise Control for automated monitoring and OAuth proxy for secure access
- Mentored a team of 2 summer interns to develop a Kubernetes operator that standardized provisioning for two
- Automated k8s application delivery using a Terraform IaC framework, cutting environment deployment time by over 90%
- Pioneered the adoption of real-time Flink, modernizing Spark jobs and reducing time to production by 40%
- Optimized Kafka cluster cost by implementing rack awareness, reducing inter-AZ network traffic cost by $1k/month

Software Engineer, Deloitte – Arlington, VA **September 2019 – September 2021**

- Managed development roadmap and technical implementation of a centralized logging solution, reducing environment downtime by 20% and resulting in $4M+ in cost savings over 3 years
- Lead development and optimization of an integrated streaming and data lake solution for analytic use cases, focusing on building extensible applications, modular interfaces and flexible architectures to support a multi-cloud environment
- Adapted big data streaming application to use cloud native services in 2-week pilot program, handling 15B events/day
- Built Flask REST API for submitting SQL queries to a Redis PubSub queue for execution
- Implemented and deployed solution to enhance monitoring of streaming data via InfluxDB and Grafana, setting up dashboards to create a real-time view of the system on the scale of millions of records per second
- Delivered tailored demos for developers, architects, data analysts, and client engagements leaders that highlighted the lower cost, reduced operational overhead, and increased performance of the redesigned cloud analytical environment

Software Engineering Intern, Tenable – Columbia, MD **June – August 2018**

- Developed a debug report application using Flask to reduce time to investigate bugs for the support engineers

RELEVANT PROJECTS

Kafka Connect Splunk (https://github.com/gjw13/kafka-connect-splunk) **March 2024**

- Introduced custom feature to transform Kafka topics to Splunk metric data format, actively forwarding 20 topics in production

APPLICABLE COURSES AND SKILLS

- Languages: Python, Scala, Java, Golang
- Technologies: Kubernetes, Kafka, Terraform, Spark, Flink, Github Actions, Docker, AWS, GCP, Grafana, Prometheus
- Management: Git, Jira, Agile, Scrum, Confluence, Kanban

EDUCATION

Georgetown University, Washington, DC – May 2019

Major: B.S. Computer Science, Minor: Business Administration

PROFESSIONAL CERTIFICATIONS

- AWS Solutions Architect – Associate (2021)

PUBLICATIONS

Author of *Plugged In: How AI Will Enhance an Evolving Sports World*

Author, published April 2018

- Interviewed experts, crafted a manuscript, designed a cover, gained endorsement quotes, marketed, and published the book